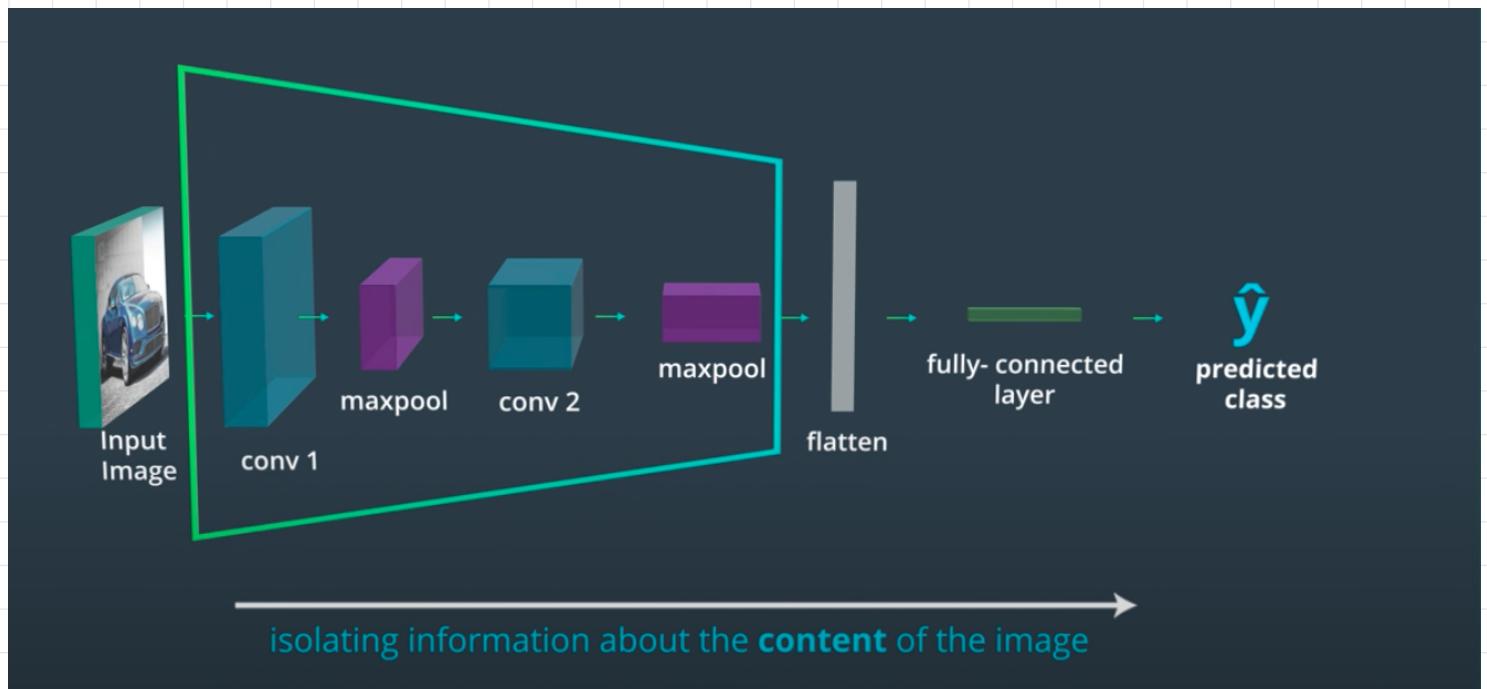


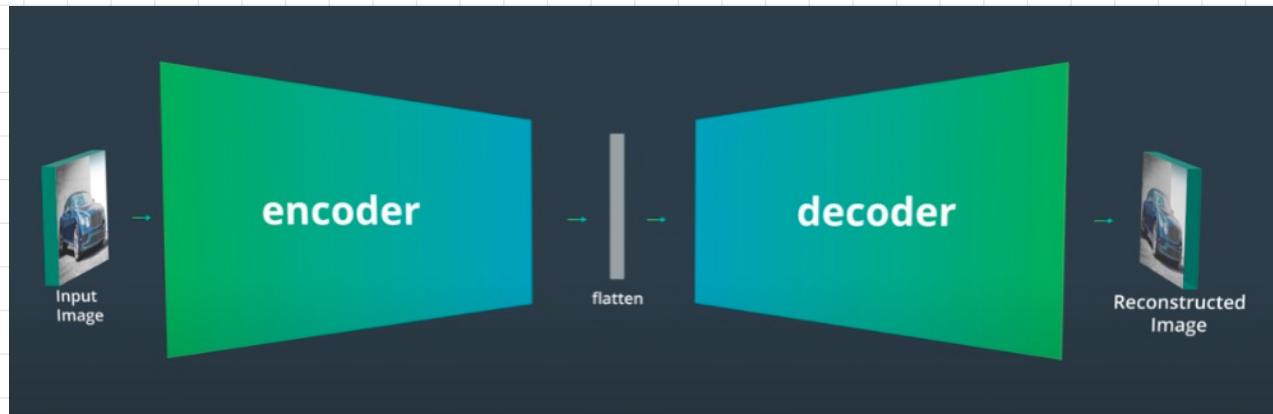
Autoencoder



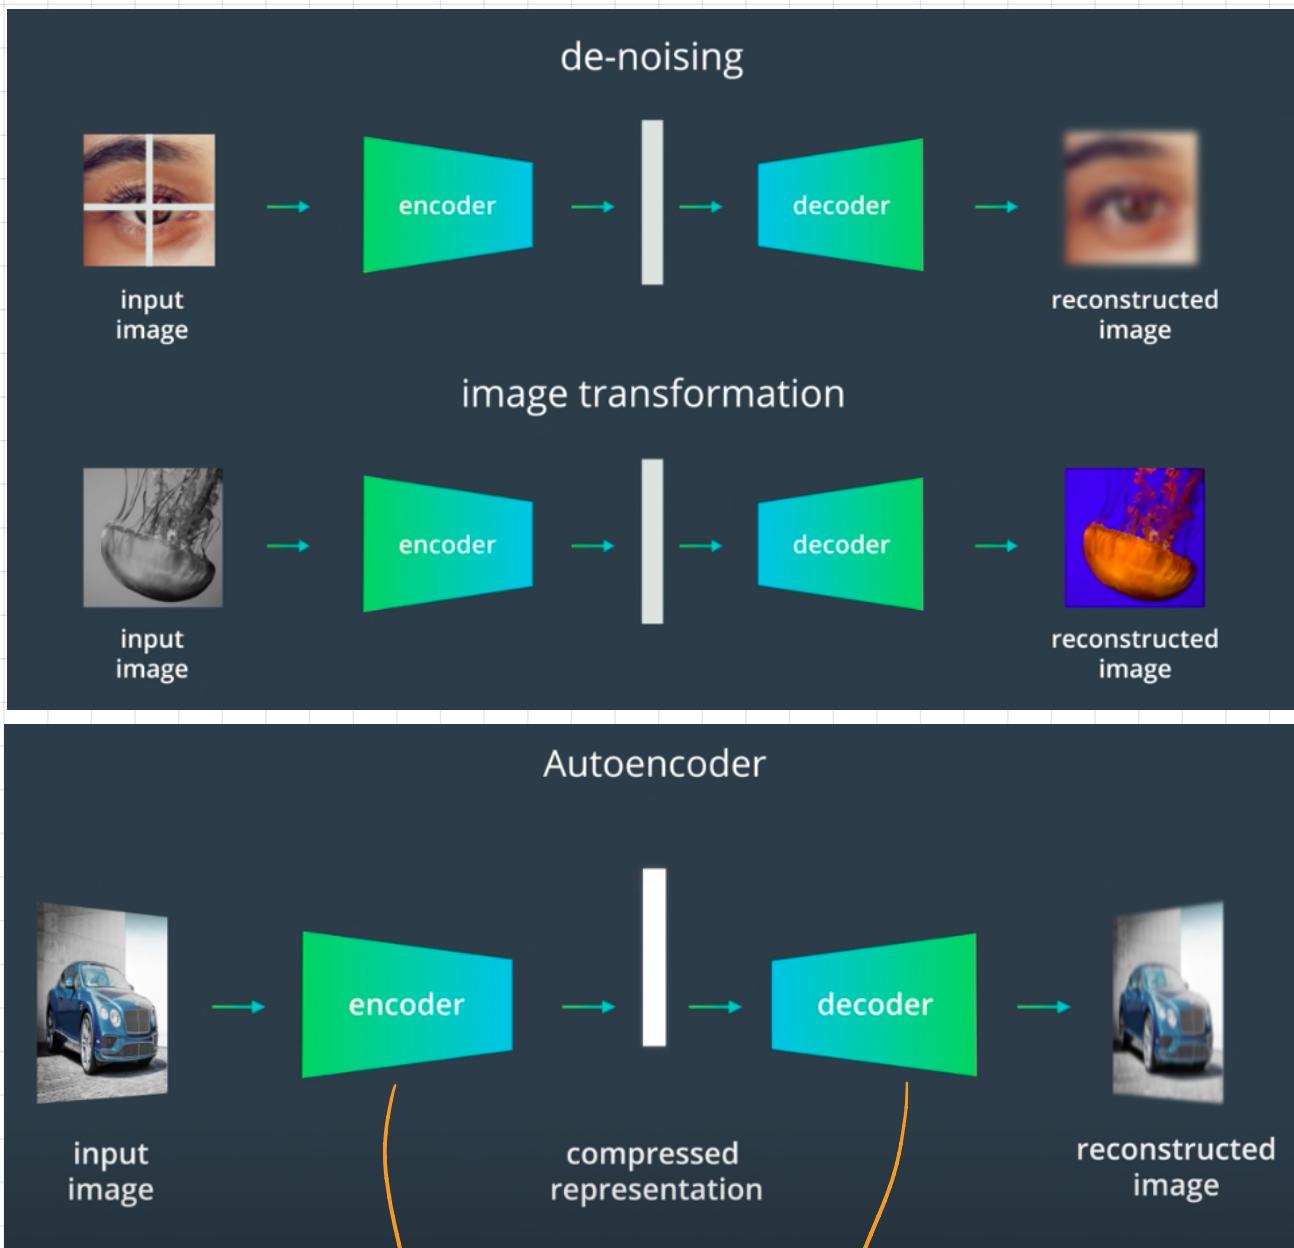
CNN Review



- Discard spatial info
- Isolate high level info about content
- Aka data compression: image compress feature vector



- Autoencoder compresses input, ↓ dimensionality
 - Read faster
- Compresses but retain content !



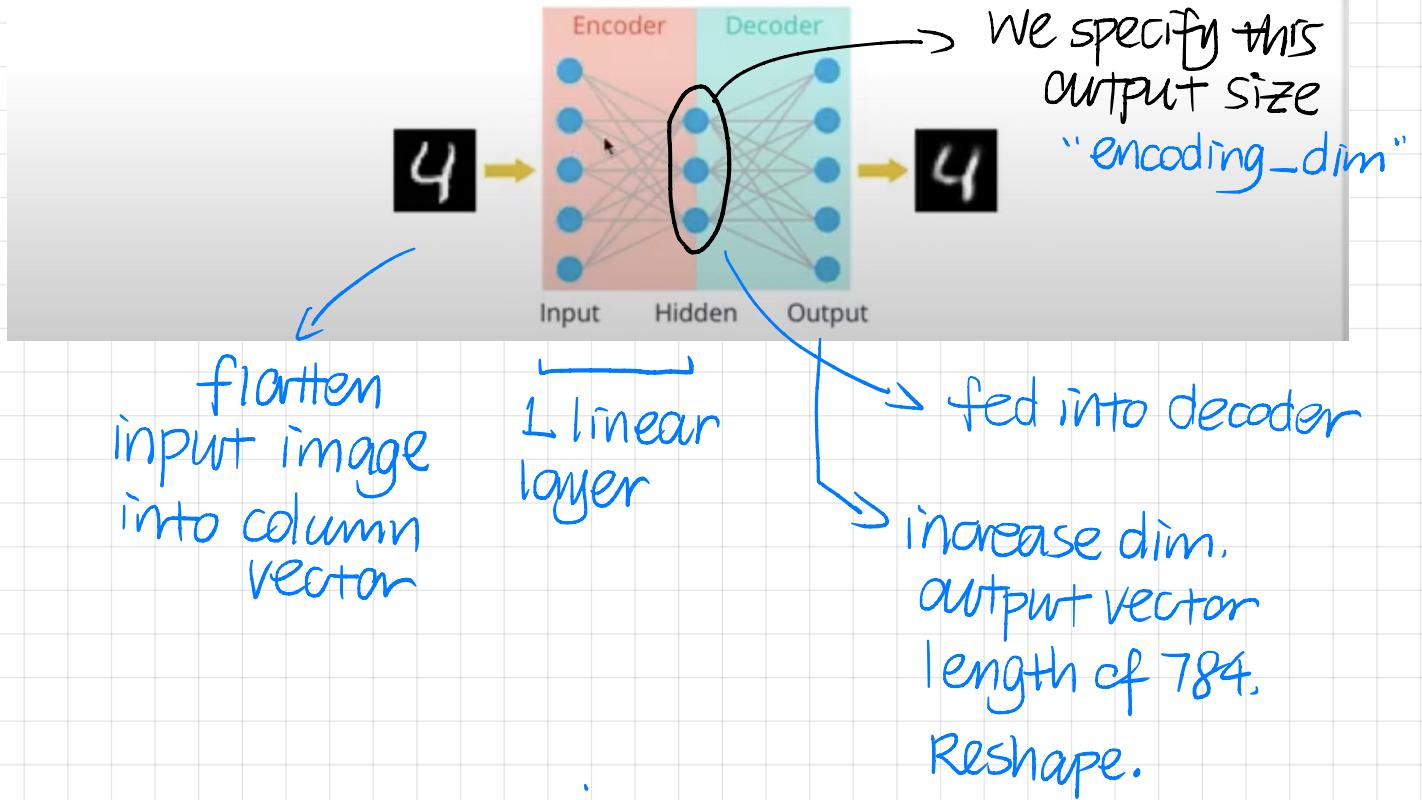
Built w/ Neural Networks

Mimimize the difference bt. input & output!

Implementation (with MLP)

We'll train an autoencoder with these images by flattening them into 784 length vectors. The images from this dataset are already normalized such that the values are between 0 and 1. Let's start by building a simple autoencoder. The encoder and decoder should be made of **one linear layer**. The units that connect the encoder and decoder will be the **compressed representation**.

Since the images are normalized between 0 and 1, we need to use a **sigmoid activation** on the output layer to get values that match this input value range.

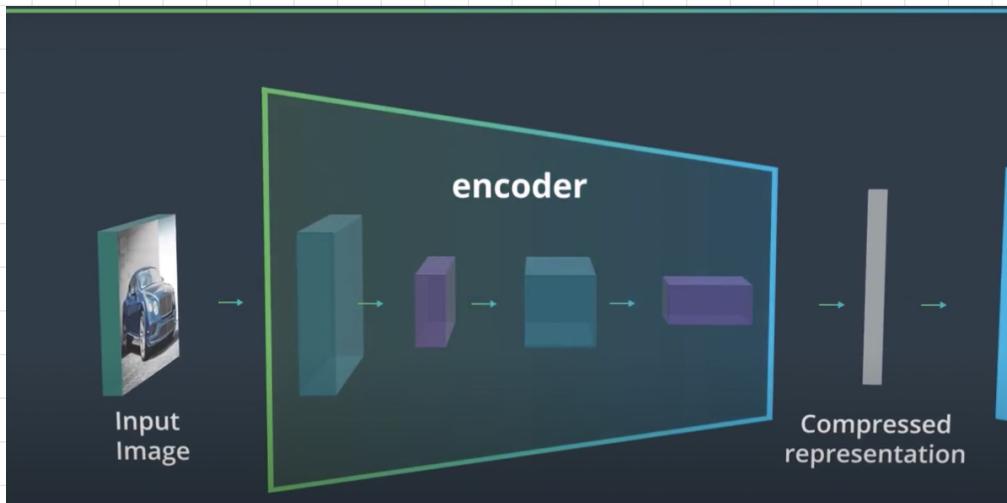


Training

- Loss: USE MSE instead of CrossEntropy
 - Comparing pixel values bt. input & output so it resembles a regression task, not a classification.
- Not interested in labels, just want to compare input & compressed images.

Implementation (With CNN)

- Advantage: Preserve content information

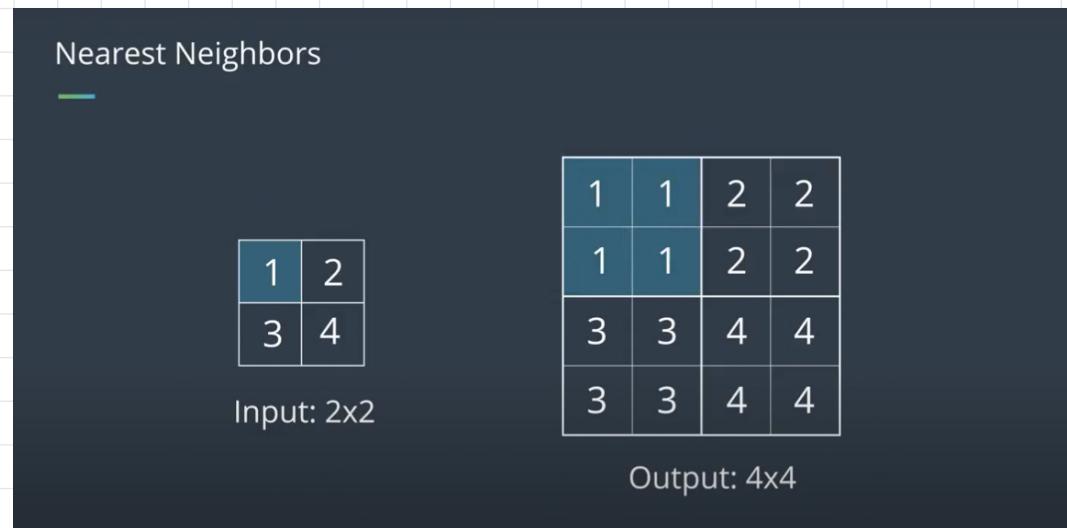


→ Spatial dims ↓

For decode, we want to increase spatial dim.

Via "un-pooling" to up-sampling

Chude.

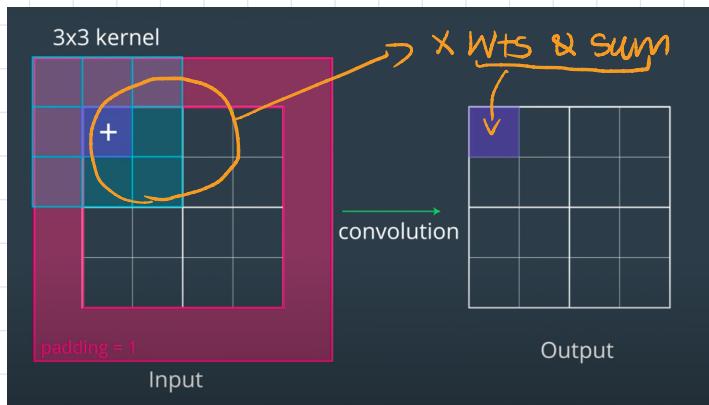


Instead, we can "learn" how to best up-sample.

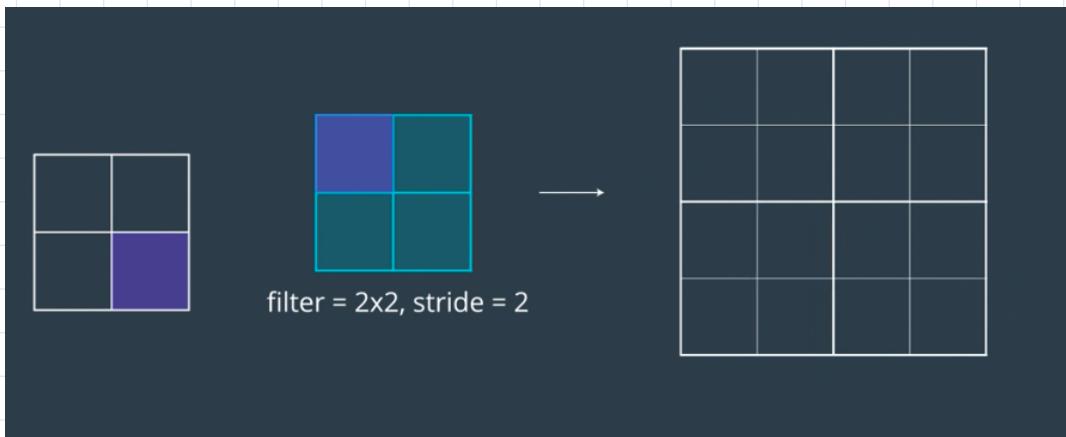
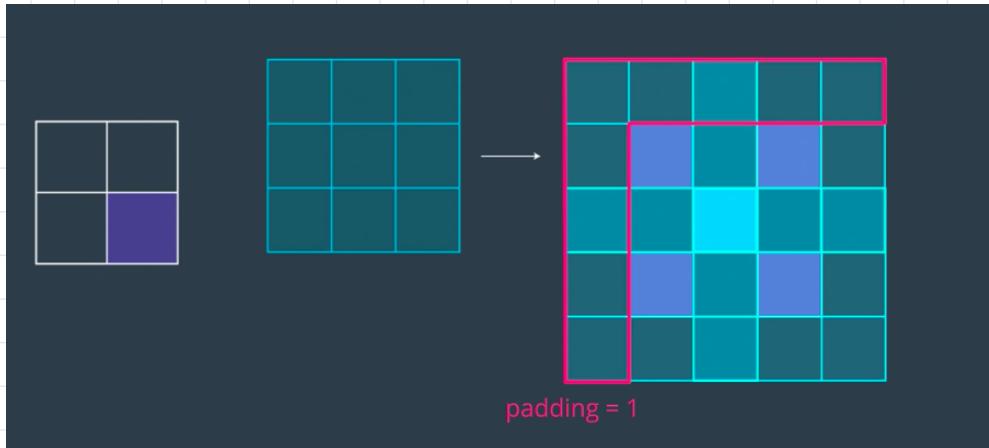
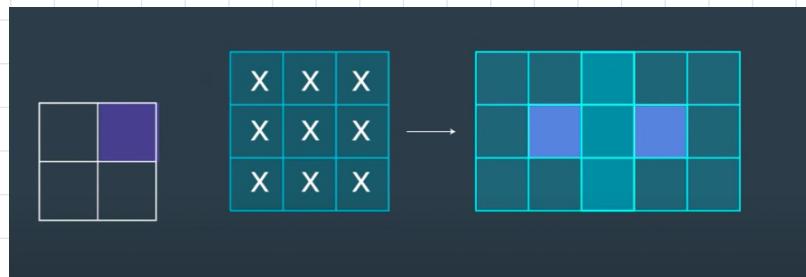
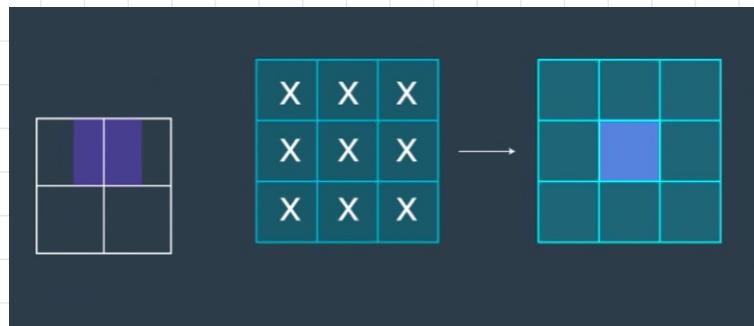
Via Transpose CNN!

Transpose Convolution

Conv. review

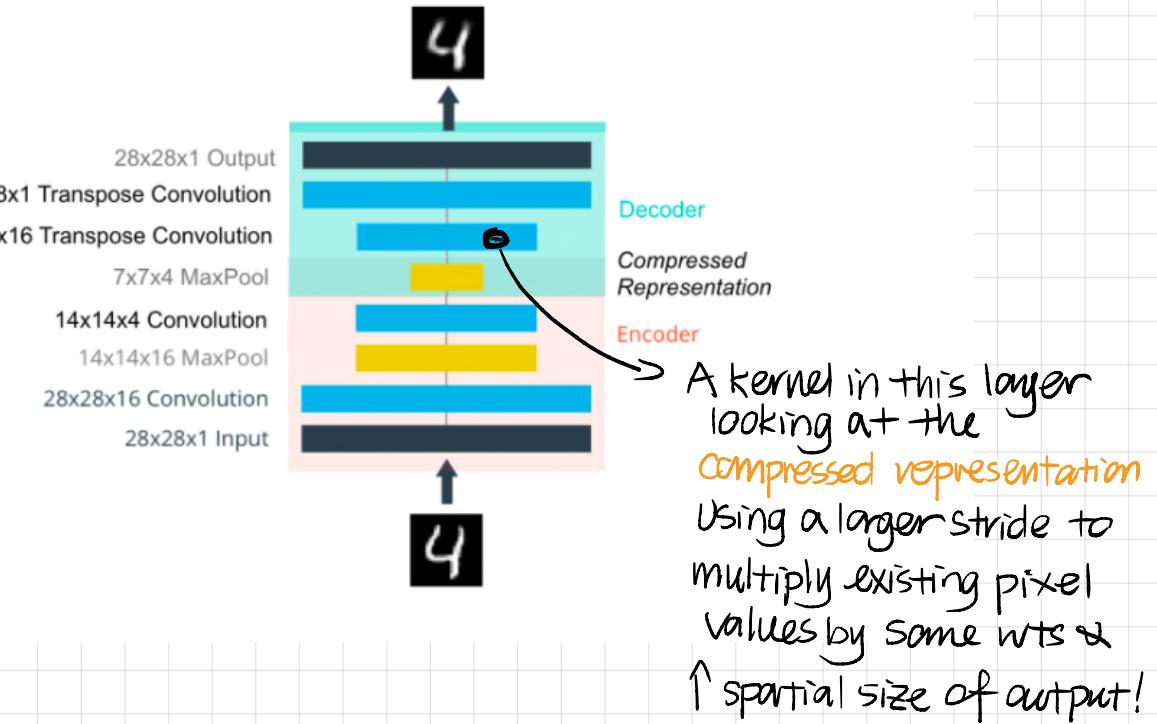


The wts here are also learned, like a regular CNN



Implementation (Method #1)

In this case, we use kernel = 2 stride = 2



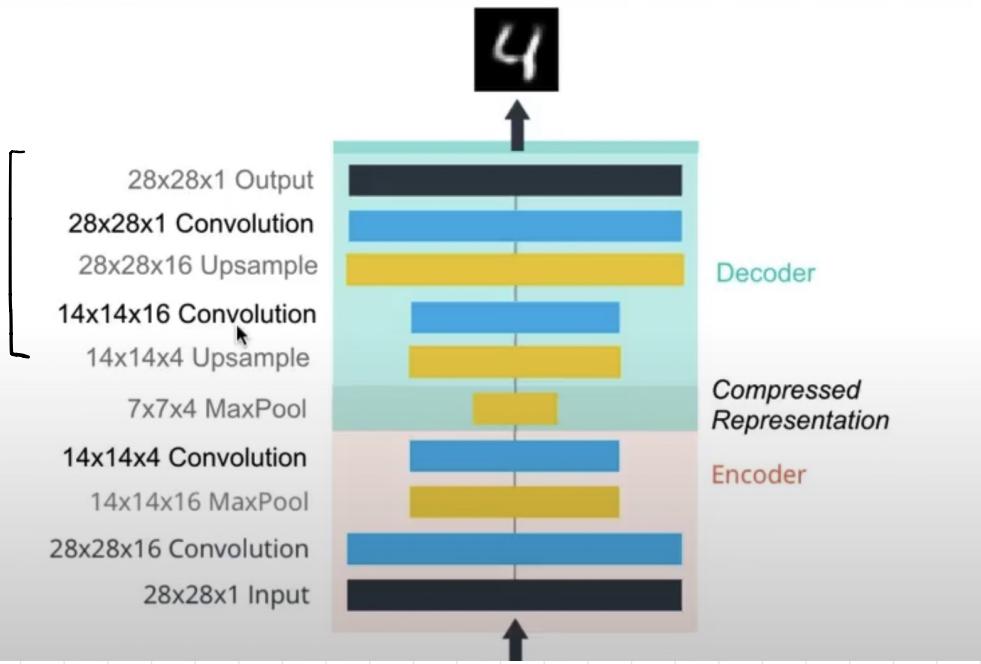
Method #2

Replace conv. T

Convolution ↑ dimension

Create filtered images
to help analyze input images

Nearest neighbor interpolation



$$\frac{(w-in-f+2p)}{s} + 1$$

$i \equiv$ width

$k \equiv$ kernel size