

Application:

Generate realistic data (Mostly for images)



Generate images
based on description
of words.

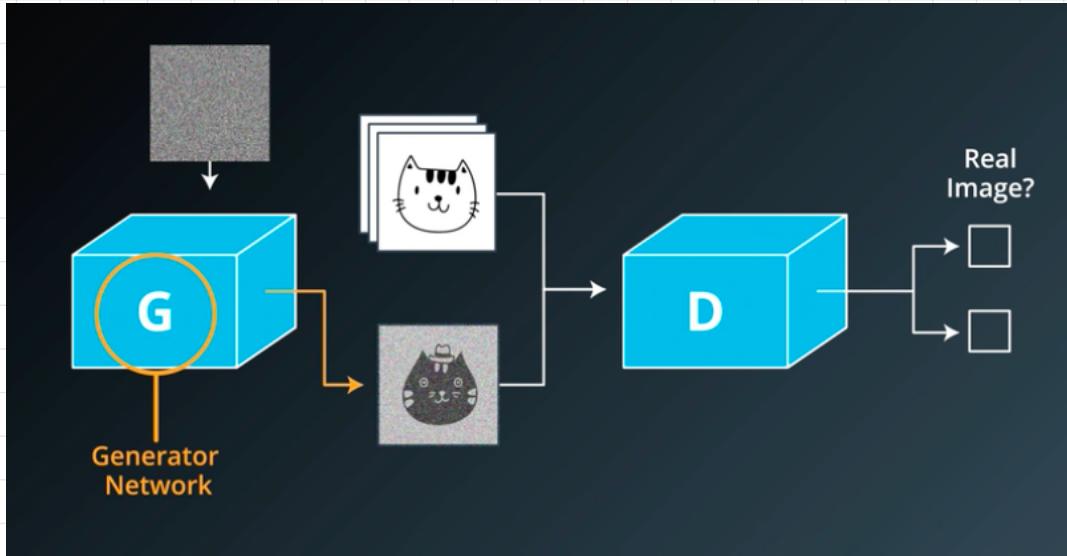
• Pix2Pix

- Drawings → real images
- Real images → cartoon characters.

• Simulated Training Set

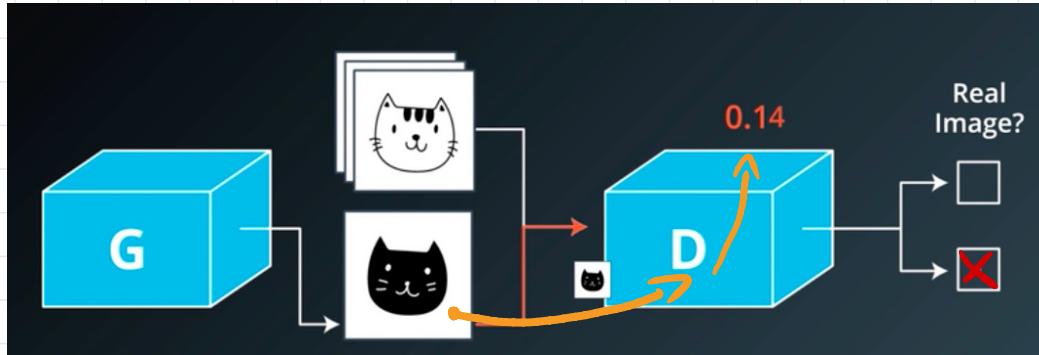
• Imitation Learning

How GANs Works?



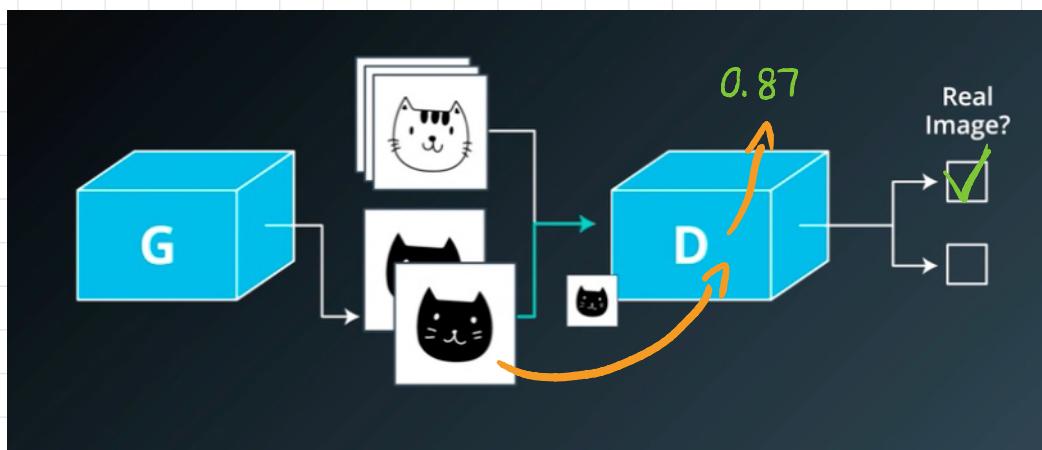
- Generator takes **Random noise** as input, runs it through differentiable functions & transform a noise into recognizable structure.
 - Output **a realistic sample**
 - A choice of noise returns diverse samples
 - **Goal:** want these generated samples to be fair samples from the distribution over real data.
 - Doesn't start out generating real/right samples.
 - **Unsupervised:** doesn't have targets associated with labels. We show GANs a bunch of samples & ask it to create more samples by **drawing from the same probability distribution**.
- Discriminator is fed with real imgs from samples & fake ones from generator 50% of the time.

The discriminator tries to:



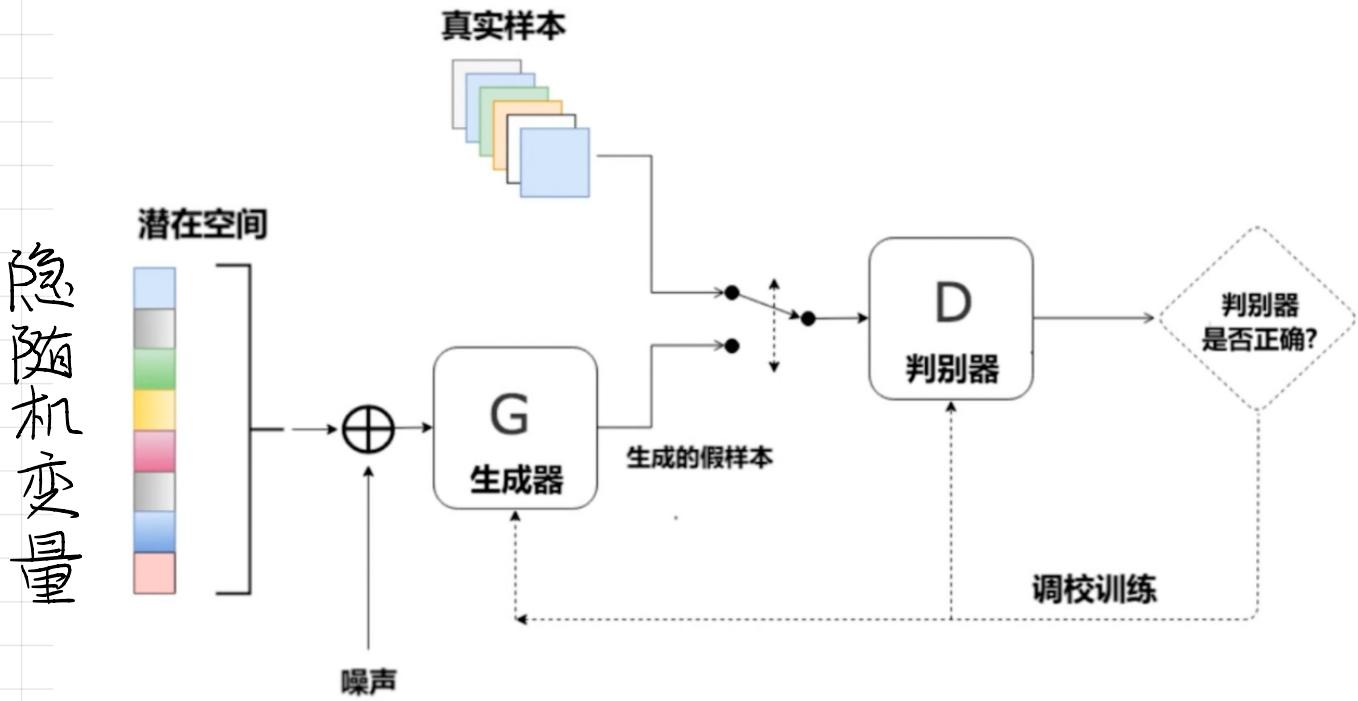
Assign a value
near 0 for
fake imgs.

The generator tries to:

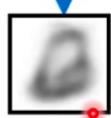
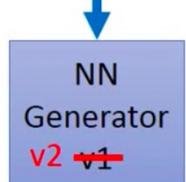


generate imgs
that will be
assigned a value
near 1.

Overtime, the generator is forced to produce near perfect samples.



Randomly sample a vector

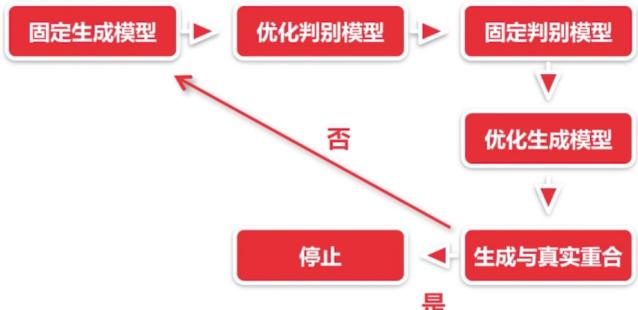


Discriminator v1

lower number
↓

Update $G(v_1)$
Use gradient descent to update parameters of G & fix D

GAN 的流程



Some Math/Theory

Maximum Likelihood Estimation

- Given data distribution $P_{\text{data}}(x)$ ↗ E.g: a joint of vectors of the pixels of an img
- We have a distribution $P_G(x; \theta)$, parameterized by θ
↓
(x & θ are similar)

E.g. A Gaussian Mixture Model, θ are means & variances of the Gaussians.

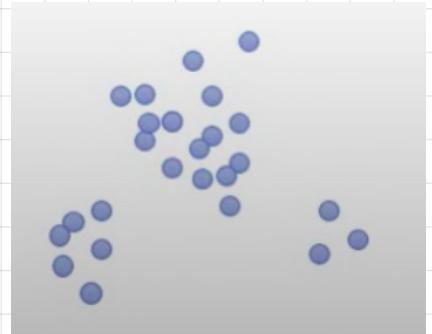
GOAL: Want to find θ such that $P_G(x; \theta) \approx P_{\text{data}}(x)$

1. Sample $\{x^1, x^2, \dots, x^m\}$ from $P_{\text{data}}(x)$

2. Compute $P_G(x^i; \theta)$

3. Likelihood of generating the models

Basically, prob of P_G generating a distribution as close to $\{x^1, \dots, x^m\}$ as possible.



$$L = \prod_{i=1}^m P_G(x^i; \theta)$$

Calculate P_G of sample $\{x^1, \dots, x^m\}$

4. Find θ^* to maximize likelihood.

$$\begin{aligned}\theta^* &= \arg \max_{\theta} \prod_{i=1}^m P_G(x^i; \theta) = \arg \max_{\theta} \log \prod_{i=1}^m P_G(x^i; \theta) \\ &= \arg \max_{\theta} \sum_{i=1}^m \log P_G(x^i; \theta)\end{aligned}$$

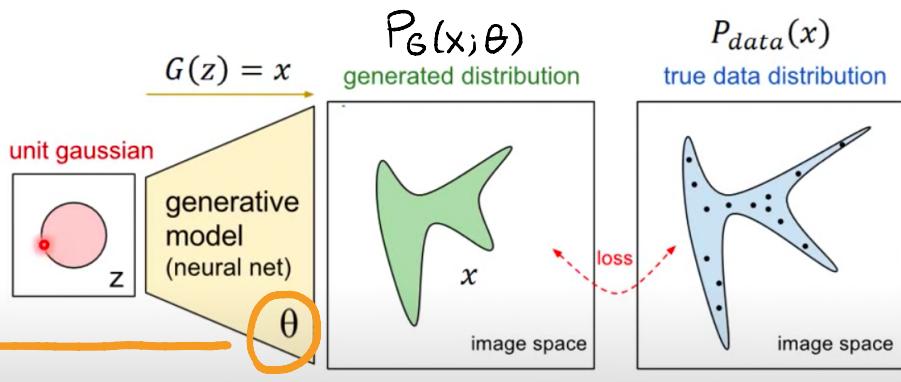
Didn't quite understand the rest...

G is determined by θ

Find parameter

θ to make G generate a distribution P_G close to $P_{\text{data}}(x)$

Now $P_G(x; \theta)$ is a NN



Z = random sampled vector from a distribution

$X = G(Z)$, also a distribution.

Prob of generating x from P_G

$$P_G(x) = \int_z P_{\text{prior}}(z) I_{[G(z)=x]} dz$$

Generator G & Discriminator D

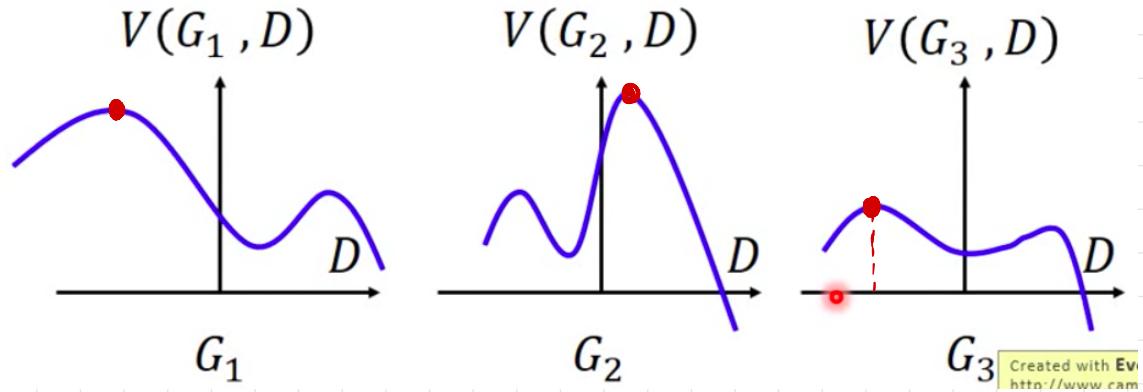
- G input z , output x
- D input x , output scalar.
 - Evaluate difference bt. $P_G(x)$ & $P_{\text{data}}(x)$
- Define a function $V(G, D)$

$$G^* = \arg \min_G \max_D V(G, D)$$

G^* = generator that generates samples near $P_{\text{data}}(x)$

Regarding Equation G*

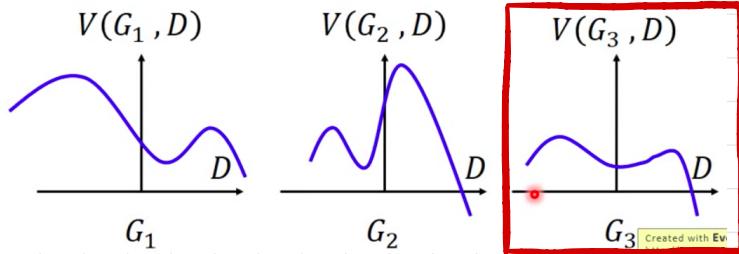
Assume that there are 3 discrete choices of G



1. Max: $G^* = \arg \min_G \max_D V(G, D)$

- Fix G, vary D
- Find a point that maximizes $V(G, D)$ (aka the red pts ↑)

2. Min $\underset{G}{\rightarrow}$



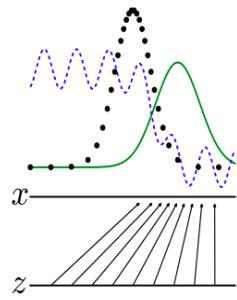
$$V = E_{x \sim P_{data}} [\log D(x)] + E_{x \sim P_G} [\log(1 - D(x))]$$

Given a generator G, $\max_D V(G, D)$ evaluate the
“difference” between P_G and P_{data}

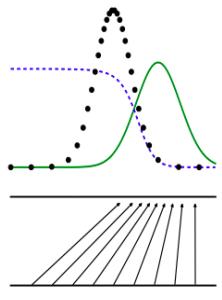


We want G_3 because $V(G, D)$ is the smallest
Therefore, difference bt. P_G & P_{data} is the smallest.

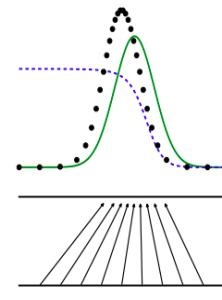
Find a G that let D on the curve $V(G, D)$ be minimal



(a)

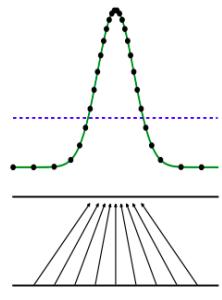


(b)



(c)

...



(d)

$\left. \begin{array}{l} D \text{ unable to} \\ \text{tell the diff.} \\ \text{bt } P_{\text{data}} \text{ &} \\ P_G, D(x) = \frac{1}{2} \end{array} \right\}$

Games & Equilibria

Game Theory:

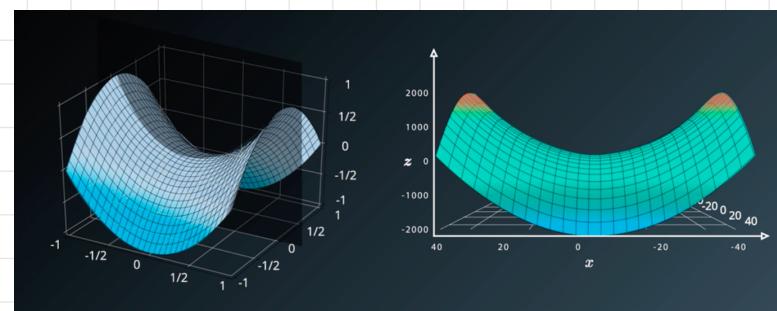
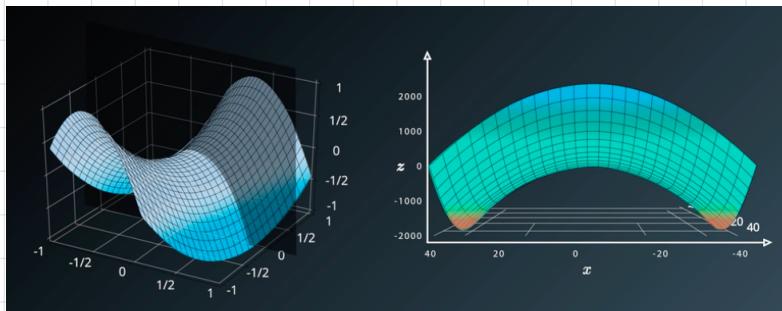
- **Equilibrium:** neither players can improve their payoff by changing their own strategy without changing the other player's strategy.
- Both sides choose their own actions **uniformly at random**

$$\text{Cost}_D = -\text{Cost}_G$$



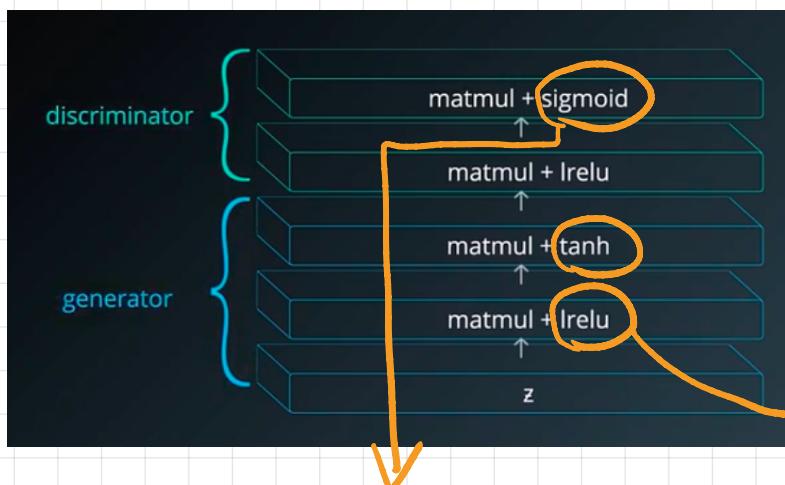
Generator wants to minimize the value function,
discriminator wants to maximize it.

∴ Equilibrium \equiv where value is max for D & min for G



Run 2 optim. algorithms, each minimizing one player's cost with respect to that player's parameters.

Tips for Training GANs

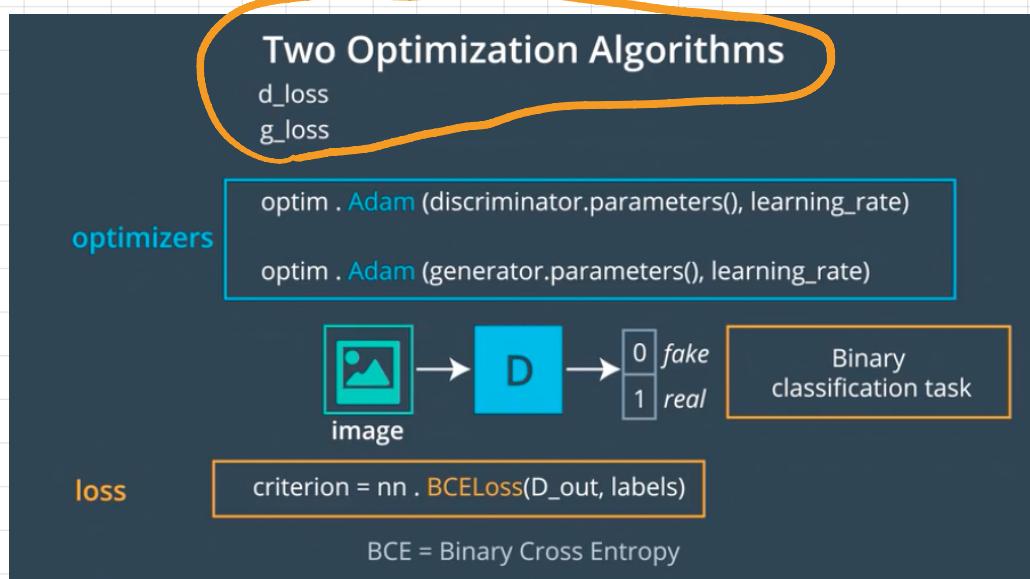


Fully Connected GANs

- Have at least 1 hidden layer for both D and G.

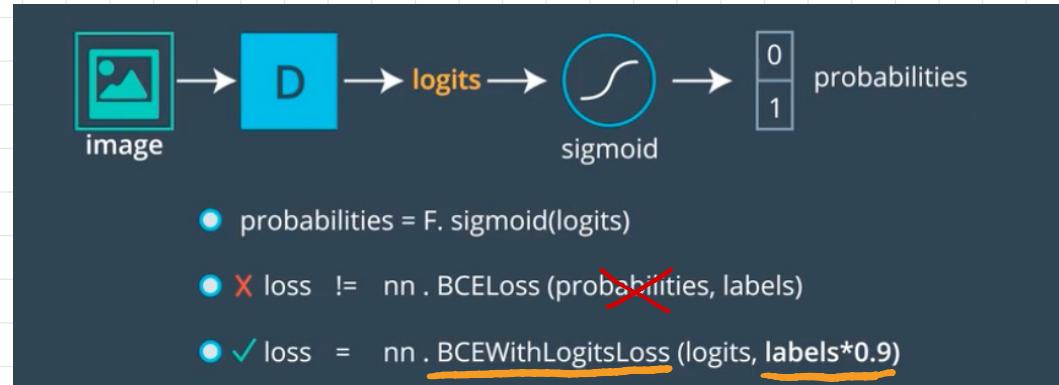
leaky ReLU for hidden units
activation function: ensure gradients to flow through the architecture

- Because the only way generator will learn is to receive a gradient from the discriminator



Discriminator Loss

Need to Use
numerically
stable cross
entropy loss
(logits)

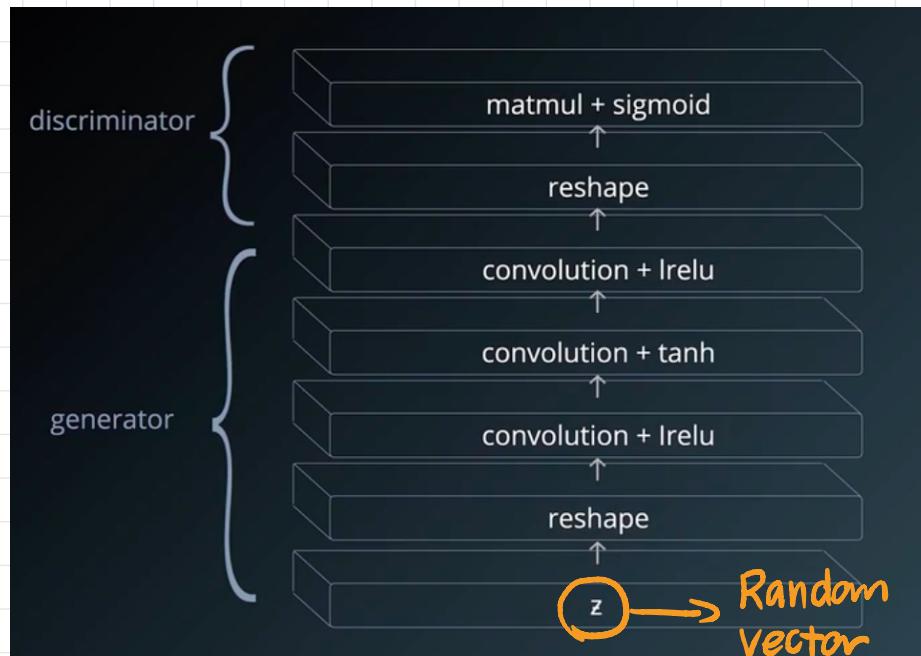


Discriminator & Generator Loss



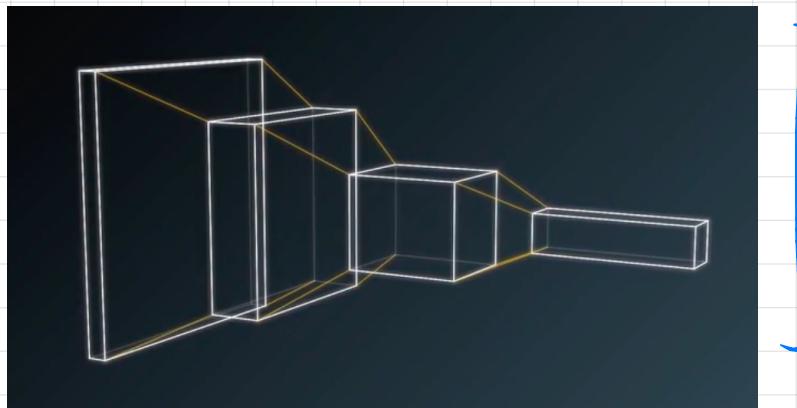
- `d_loss = nn.BCEWithLogitsLoss(logits, labels*0.9)`
- `g_loss = nn.BCEWithLogitsLoss(logits, flipped_labels)`

Convolutional GANs



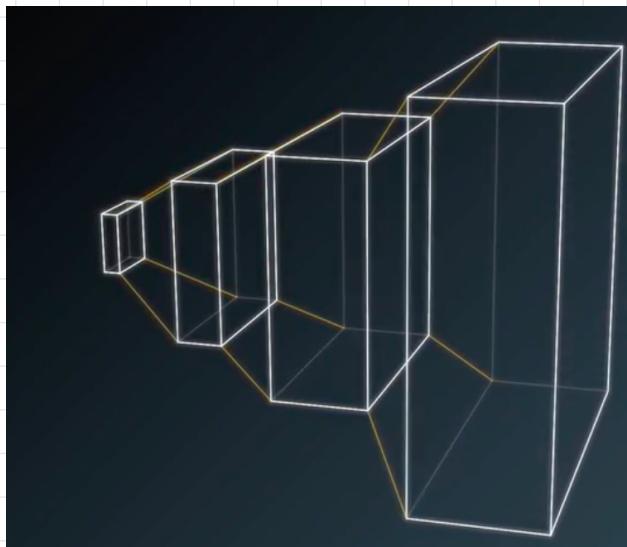
- A minibatch of z forms a matrix.
- A convolution expects a 4D tensor for input:
 - 1 axis for different examples in minibatch
 - 1 axis for diff. feature maps
 - 2 axes for width & height

Opposite to a classifier conv net:



Height ↓
width ↓
depth ↑

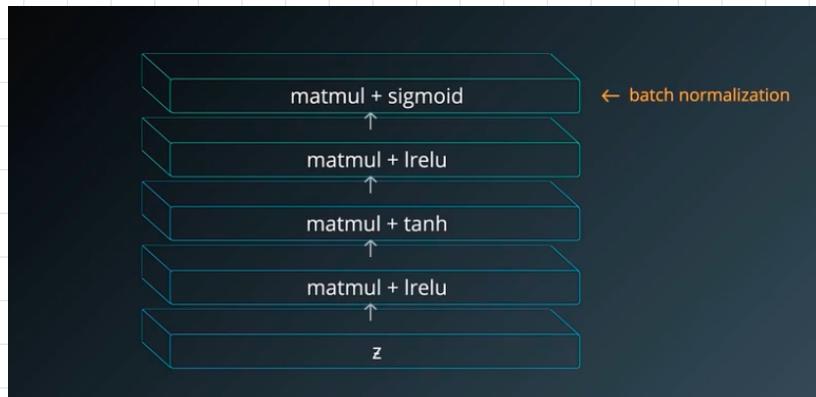
When using conv net as a generator net:



Do the opposite!

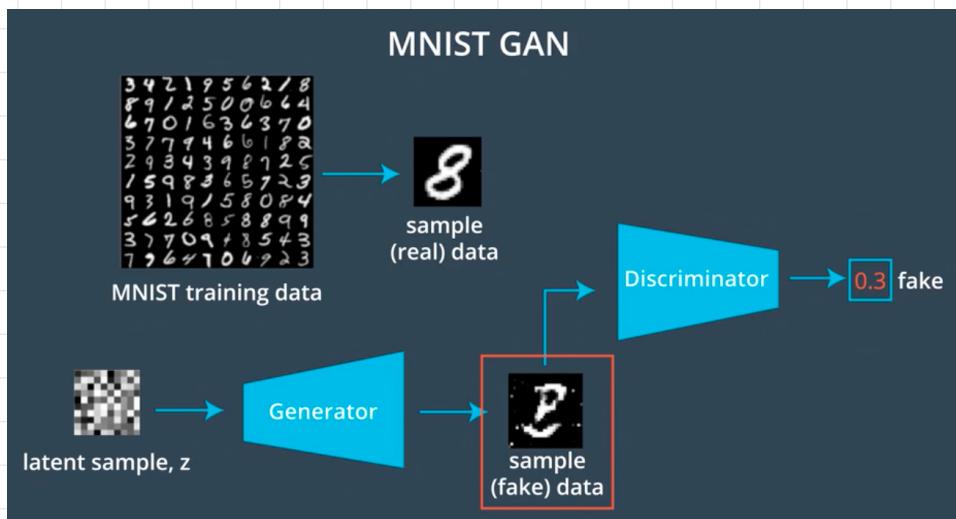
DCGANs.

Batch Norm



Batch norm on
all layers but
input & final
classifier

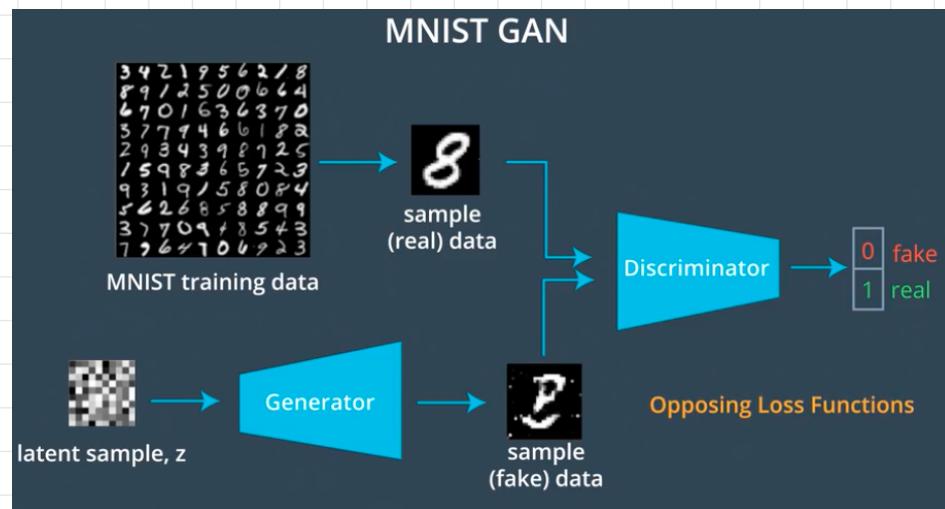
A Final Overview



Generator is learning & getting better at creating better samples & fool D to mis-labeled fake data as true.

Goal:

Maximize the error
Made by D.



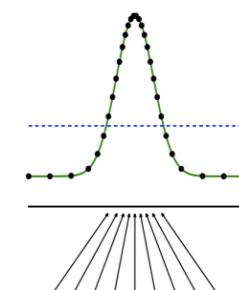
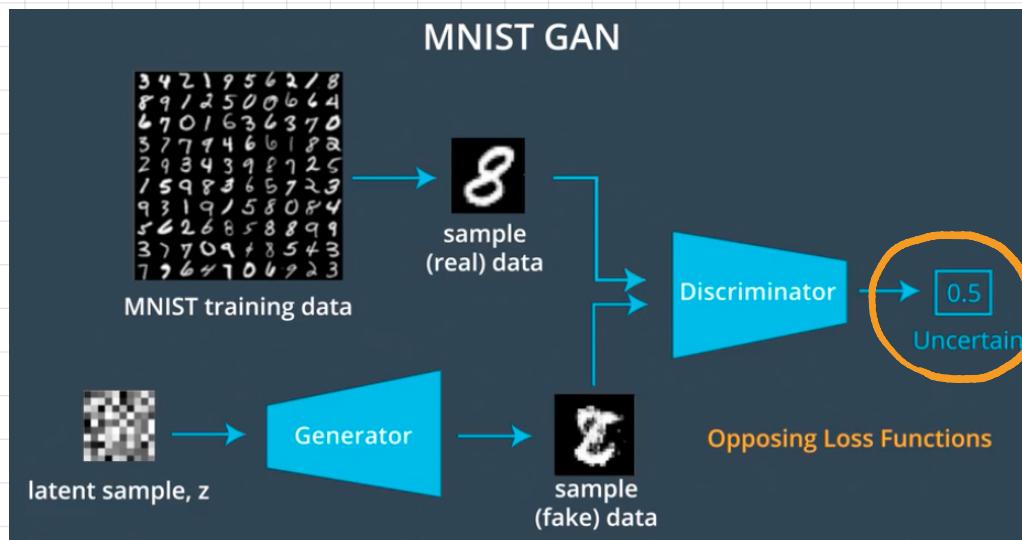
D is also training, by being fed w/ alternating real & fake imgs, getting better at identifying.

Goal:

Minimize error rate.

Recall equilibrium & Game Theory:

Two players at equilibrium, no move tradeoff



(d)

$$D(x) = \frac{1}{2}$$

